

# Lecture 4: Structured Prediction Models

Kai-Wei Chang  
CS @ UCLA  
[kw@kwchang.net](mailto:kw@kwchang.net)

Couse webpage: <https://uclanlp.github.io/CS269-17/>

# Previous Lecture

- ❖ Binary linear classification

- ❖ Perceptron, SVMs, Logistic regression, Naïve Bayes

- ❖ **Output:**  $y \in \{1, -1\}$

- ❖ Multi-class classification

- ❖ Multiclass Perceptron, Multiclass SVM...

- ❖ **Output:**  $y \in \{1, 2, 3, \dots K\}$

# What we have seen: multiclass

- ❖ Reducing multiclass to binary
  - ❖ One-against-all & One-vs-one
  - ❖ Error correcting codes
  - ❖ Extension: Learning to search
- ❖ Training a single classifier
  - ❖ Multiclass Perceptron: Kesler's construction
  - ❖ Multiclass SVMs: Crammer&Singer formulation
  - ❖ Multinomial logistic regression
  - ❖ Extension: Graphical models

# This lecture


- ❖ What is structured output?
- ❖ Multiclass as structure
- ❖ Sequence as structure
- ❖ General graph structure

# Global decisions

- ❖ “Understanding” is a global decision
  - ❖ Several local decisions play a role
  - ❖ There are mutual dependencies on their outcome.
- ❖ Essential to make coherent decisions
  - ❖ Joint, Global Inference

# Inference with Constraints

[Roth&Yih'04,07,.....]



Models could be learned separately/jointly; constraints may come up only at decision time.

# Inference with Constraints [Roth&Yih'04,07,...]

other	0.05
per	0.85
loc	0.10

other	0.10
per	0.60
loc	0.30

other	0.05
per	0.50
loc	0.45

Bernie's wife, Jane, is a native of Brooklyn

$E_1$

$E_2$

$E_3$

$R_{12}$

$R_{23}$

irrelevant	0.05
spouse_of	0.45
born_in	0.50

irrelevant	0.10
spouse_of	0.05
born_in	0.85


Models could be learned separately/jointly; constraints may come up only at decision time.

# Inference with Constraints [Roth&Yih'04,07,...]

other	0.05
<b>per</b>	<b>0.85</b>
loc	0.10

other	0.10
<b>per</b>	<b>0.60</b>
loc	0.30

other	0.05
per	0.50
<b>loc</b>	<b>0.45</b>



irrelevant	0.05
<b>spouse_of</b>	<b>0.45</b>
born_in	0.50

irrelevant	0.10
spouse_of	0.05
<b>born_in</b>	<b>0.85</b>

Models could be learned separately/jointly; constraints may come up only at decision time.




# Structured output is...

❖ A predefined structure

Predicate	A0	A1	Location
Head	The bus	Nairobi in Kenya	-

❖ Can be represented as a graph



# Sequential tagging

- ❖ The process of assigning a part-of-speech to each word in a collection (sentence).

WORD

tag

**the**

**DET**

**koala**

**N**

**put**

**V**

**the**

**DET**

**keys**

**N**

**on**

**P**

**the**

**DET**

**table**

**N**

# Let's try

Don't worry! There is no problem with your eyes or computer.

෧/DT ෧෪/NN ෦ 10/VBZ ෨෧෧ 5/VBG ෧/DT  
෨෧0/NN .

෧/DT ෧෪ 4/NN ෦ 10/VBZ 9 1 5 5 0 5/VBG .

෧/DT ෧෪ 5/NN ෦ 10/VBZ 10 0 5/VBG ෦ 5/VBG .

෧/DT ෧෧෧ 7 7 5/JJ ෧෦ 9/NN

What is the POS tag sequence of the following sentence?

෧ ෧෧෧ 7 7 5 ෨෧෦ 3 ෧ 10 10 0 5/VBG ෦ 5/VBG .

# Let's try

❖ **෧/DT ෧෦/NN ෦෦/VBZ ෧෦෦෦/VBG ෧/DT ෧෦/NN** /.

a/DT dog/NN is/VBZ chasing/VBG a/DT cat/NN ./.

❖ **෧/DT ෧෦/NN ෦෦/VBZ ෦෦෦/VBG** /.

a/DT fox/NN is/VBZ running/VBG ./.

❖ **෧/DT ෧෦/NN ෦෦/VBZ ෦෦෦/VBG** /.

a/DT boy/NN is/VBZ singing/VBG ./.

❖ **෧/DT ෦෦෦/JJ ෧෦/NN**

a/DT happy/JJ bird/NN

❖ **෧ ෦෦෦෦ ෧෦෦ ෧෦෦ ෦෦෦෦ ෦෦෦෦** 

a happy cat was singing .

# How you predict the tags?

- ❖ Two types of information are useful
  - ❖ Relations between words and tags
  - ❖ Relations between tags and tags
    - ❖ DT NN, DT JJ NN...
    - ❖ Fed in “The Fed” is a Noun because it follows a Determiner

# Combinatorial optimization problem

$$\hat{y} = \operatorname{argmax}_{y \in \mathcal{Y}} f(y; \mathbf{w}, \mathbf{x})$$

input

model parameters

output space

❖ Inference/Test: given  $\mathbf{w}, \mathbf{x}$ , solve  $\operatorname{argmax}$


❖ Learning/Training: find a good  $\mathbf{w}$

# Challenges with structured output

- ❖ We cannot train a separate weight vector for each possible inference outcome (why?)
  - ❖ For multi-class we train one weight vector for each class
- ❖ We cannot enumerate all possible structures for inference
  - ❖ Inference for multiclass was easy

# Deal with combinatorial output

- ❖ Decompose the output into **parts** that are labeled
- ❖ Define a **graph** to represent
  - ❖ how the parts **interact** with each other
  - ❖ These labeled interacting parts are **scored**; the total score for the graph is the sum of scores of each part
  - ❖ an **inference algorithm** to assign labels to all the parts





# A history-based model

$$P(x_1, x_2, \dots, x_n) = \prod_{i=1}^n P(x_i | x_1, x_2, \dots, x_{i-1})$$

- ❖ Each token is dependent on all the tokens that came before it
- ❖ Simple conditioning
- ❖ Each  $P(x_i | \dots)$  is a multinomial probability distribution over the tokens



# Example: A Language model


It was a bright cold day in April.

$$P(\text{It was a bright cold day in April}) =$$

$P(\text{It}) \times$   Probability of a word starting a sentence

$P(\text{was}|\text{It}) \times$   Probability of a word following “It”

$P(\text{a}|\text{It was}) \times$   Probability of a word following “It was”

$P(\text{bright}|\text{It was a}) \times$   Probability of a word following “It was a”

$P(\text{cold}|\text{It was a bright}) \times$

$P(\text{day}|\text{It was a bright cold}) \times \dots$

# A history-based model

$$P(x_1, x_2, \dots, x_n) = \prod_{i=1}^n P(x_i | x_1, x_2, \dots, x_{i-1})$$

- ❖ Each token is dependent on all the tokens that came before it
  - ❖ Simple conditioning
  - ❖ Each  $P(x_i | \dots)$  is a multinomial probability distribution over the tokens
- ❖ What is the problem here?
  - ❖ How many parameters do we have?
    - ❖ Grows with the size of the sequence!

# Solution: Lose the history

## Discrete Markov Process

- ❖ A system can be in one of  $K$  states at a time
- ❖ State at time  $t$  is  $x_t$

- ❖ **First-order Markov assumption**

The state of the system at any time is ***independent*** of the full sequence history given the previous state

$$P(x_1, x_2, \dots, x_n) = \prod_{i=1}^n P(x_i | x_{i-1})$$


- ❖ Defined by two sets of probabilities:
  - ❖ **Initial** state distribution:  $P(x_1 = S_j)$
  - ❖ State **transition** probabilities:  $P(x_i = S_j \mid x_{i-1} = S_k)$

# Example: Another language model

It was a bright cold day in April

$P(\text{It was a bright cold day in April}) =$

$P(\text{It}) \times$   Probability of a word starting a sentence

$P(\text{was}|\text{It}) \times$   Probability of a word following “It”

$P(\text{a}|\text{was}) \times$   Probability of a word following “was”

$P(\text{bright}|\text{a}) \times$   Probability of a word following “a”


$P(\text{cold}|\text{bright}) \times$

$P(\text{day}|\text{cold}) \times \dots$

If there are  $K$  tokens/states, how many parameters do we need?  $O(K^2)$

# Example: The weather

- ❖ Three states: rain, cloudy, sunny



- ❖ Observations are Markov chains:

Eg: *cloudy sunny sunny rain*

Probability of the sequence =


$$P(\text{cloudy}) P(\text{sunny}|\text{cloudy}) P(\text{sunny} | \text{sunny}) P(\text{rain} | \text{sunny})$$

Initial probability

Transition probabilities

# Example: The weather

- ❖ Three states: rain, cloudy, sunny



- ❖ Observe: These probabilities define the model;  
can find  $P(\text{any sequence})$

Eg: cloudy, sunny, sunny, rain

Probability of the sequence =

$$P(\text{cloudy}) P(\text{sunny}|\text{cloudy}) P(\text{sunny} | \text{sunny}) P(\text{rain} | \text{sunny})$$

Initial probability

Transition probabilities

# Outline

- ❖ Sequence models
- ❖ *Hidden Markov models*
  - ❖ Inference with HMM
  - ❖ Learning
- ❖ Conditional Models and Local Classifiers
- ❖ Global models
  - ❖ Conditional Random Fields
  - ❖ Structured Perceptron for sequences



# Hidden Markov Model

- ❖ Discrete Markov Model:


- ❖ States follow a Markov chain
- ❖ *Each **state** is an observation*

- ❖ Hidden Markov Model:

- ❖ States follow a Markov chain
- ❖ **States are not observed**
- ❖ Each state stochastically emits an observation

# Toy part-of-speech example

## The Fed raises interest rates



# Joint model over states and observations

## ❖ Notation

- ❖ Number of states =  $K$ , Number of observations =  $M$
- ❖  $\pi$ : Initial probability over states ( $K$  dimensional vector)
- ❖  $A$ : Transition probabilities ( $K \times K$  matrix)
- ❖  $B$ : Emission probabilities ( $K \times M$  matrix)

## ❖ Probability of states and observations

- ❖ Denote states by  $y_1, y_2, \dots$  and observations by  $x_1, x_2, \dots$

$$\begin{aligned} P(x_1, x_2, \dots, x_n, y_1, y_2, \dots, y_n) &= P(y_1) \prod_{i=1}^{n-1} P(y_{i+1}|y_i) \prod_{i=1}^n P(x_i|y_i) \\ &= \pi_{y_1} \prod_{i=1}^{n-1} A_{y_i, y_{i+1}} \prod_{i=1}^n B_{y_i, x_i} \end{aligned}$$

# Other applications

- ❖ Speech recognition
  - ❖ Input: Speech signal
  - ❖ Output: Sequence of words
- ❖ NLP applications
  - ❖ Information extraction
  - ❖ Text chunking
- ❖ Computational biology
  - ❖ Aligning protein sequences
  - ❖ Labeling nucleotides in a sequence as exons, introns, etc.

# Three questions for HMMs

[Rabiner 1999]

1. Given an observation sequence,  $x_1, x_2, \dots, x_n$  and a model  $(\pi, A, B)$ , how to efficiently calculate the probability of the observation?
2. Given an observation sequence,  $x_1, x_2, \dots, x_n$  and a model  $(\pi, A, B)$ , how to efficiently calculate the most probable state sequence?

Inference

3. How to calculate  $(\pi, A, B)$  from observations?

Learning

# Outline

- ❖ Sequence models
- ❖ Hidden Markov models
  - ❖ *Inference with HMM*
  - ❖ Learning
- ❖ Conditional Models and Local Classifiers
- ❖ Global models
  - ❖ Conditional Random Fields
  - ❖ Structured Perceptron for sequences

# Most likely state sequence

- ❖ Input:

- ❖ A hidden Markov model  $(\pi, A, B)$
- ❖ An observation sequence  $\mathbf{x} = (x_1, x_2, \dots, x_n)$

- ❖ Output: A state sequence  $\mathbf{y} = (y_1, y_2, \dots, y_n)$  that corresponds to

- ❖ *Maximum a posteriori* inference (MAP inference)

$$\arg \max_{\mathbf{y}} P(\mathbf{y} | \mathbf{x}, \pi, A, B)$$

- ❖ Computationally: combinatorial optimization

# MAP inference

❖ We want  $\arg \max_{\mathbf{y}} P(\mathbf{y}|\mathbf{x}, \pi, A, B)$

❖ We have defined

$$P(x_1, x_2, \dots, x_n, y_1, y_2, \dots, y_n) = P(y_1) \prod_{i=1}^{n-1} P(y_{i+1}|y_i) \prod_{i=1}^n P(x_i|y_i)$$

❖ But  $P(\mathbf{y}|\mathbf{x}, \pi, A, B) \propto P(\mathbf{x}, \mathbf{y}|\pi, A, B)$

❖ And we don't care about  $P(\mathbf{x})$  we are maximizing over  $\mathbf{y}$

❖ So,  $\arg \max_{\mathbf{y}} P(\mathbf{y}|\mathbf{x}, \pi, A, B) = \arg \max_{\mathbf{y}} P(\mathbf{y}, \mathbf{x}|\pi, A, B)$



# How many possible sequences?

The Fed raises interest rates

Determiner	Verb	Verb	Verb	Verb
	Noun	Noun	Noun	Noun

1	2	2	2	2
---	---	---	---	---

List of allowed tags for each word

In this simple case, 16 sequences ( $1 \times 2 \times 2 \times 2 \times 2$ )

# Naïve approaches

## 1. Try out every sequence

- ❖ Score the sequence  $\mathbf{y}$  as  $P(\mathbf{y}|\mathbf{x}, \pi, A, B)$
- ❖ Return the highest scoring one
- ❖ What is the problem?
  - ❖ Correct, but slow,  $O(K^n)$

## 2. Greedy search

- ❖ Construct the output left to right
- ❖ For each  $i$ , elect the best  $y_i$  using  $y_{i-1}$  and  $x_i$
- ❖ What is the problem?
  - ❖ Incorrect but fast,  $O(nK)$

**Solution:** Use the independence assumptions

**Recall:** The first order Markov assumption

The state at token  $i$  is only influenced by the previous state, the next state and the token itself

Given the adjacent labels, the others do not matter


Suggests a recursive algorithm

# Jason's ice cream

#cones

	$p(\dots C)$	$p(\dots H)$	$p(\dots \text{START})$
(1 ...)	<b>0.5</b>	<b>0.1</b>	
(2 ...)	<b>0.4</b>	<b>0.2</b>	
(3 ...)	<b>0.1</b>	<b>0.7</b>	
(C ...)	<b>0.8</b>	<b>0.2</b>	<b>0.5</b>
(H ...)	<b>0.2</b>	<b>0.8</b>	<b>0.5</b>

❖ Best tag sequence for  $P("1,2,1")$ ?



# Deriving the recursive algorithm


$$P(x_1, x_2, \dots, x_n, y_1, y_2, \dots, y_n) = P(y_1) \prod_{i=1}^{n-1} P(y_{i+1}|y_i) \prod_{i=1}^n P(x_i|y_i)$$

$$\max_{y_1, y_2, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots P(y_2|y_1)P(x_2|y_2)P(y_1)P(x_1|y_1)$$

Transition probabilities

Emission probabilities

Initial probability




# Deriving the recursive algorithm

$$P(x_1, x_2, \dots, x_n, y_1, y_2, \dots, y_n) = P(y_1) \prod_{i=1}^{n-1} P(y_{i+1}|y_i) \prod_{i=1}^n P(x_i|y_i)$$

$$\begin{aligned} & \max_{y_1, y_2, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots P(y_2|y_1)P(x_2|y_2)P(y_1)P(x_1|y_1) \\ &= \max_{y_2, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots \max_{y_1} P(y_2|y_1)P(x_2|y_2)P(y_1)P(x_1|y_1) \end{aligned}$$

The only terms that depend on  $y_1$



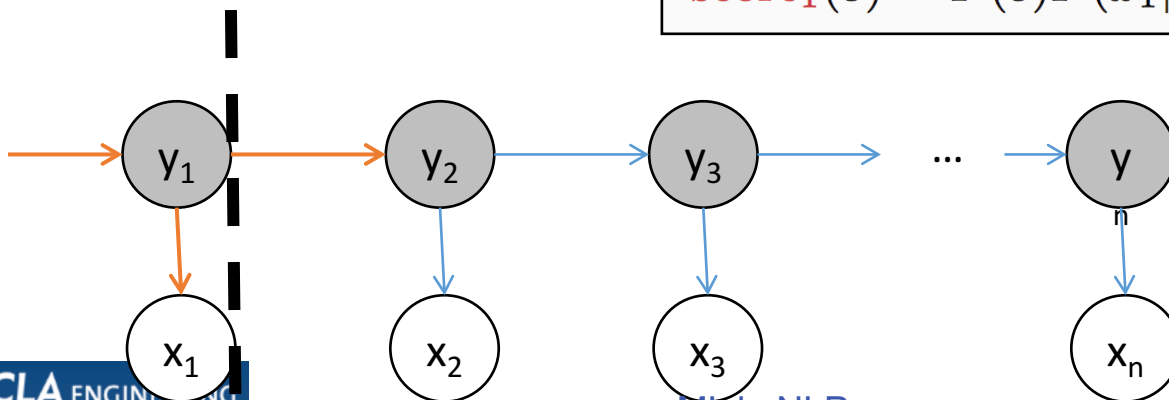
# Deriving the recursive algorithm

$$P(x_1, x_2, \dots, x_n, y_1, y_2, \dots, y_n) = P(y_1) \prod_{i=1}^{n-1} P(y_{i+1}|y_i) \prod_{i=1}^n P(x_i|y_i)$$

$$\begin{aligned} & \max_{y_1, y_2, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots P(y_2|y_1)P(x_2|y_2)P(y_1)P(x_1|y_1) \\ &= \max_{y_2, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots \max_{y_1} P(y_2|y_1)P(x_2|y_2)P(y_1)P(x_1|y_1) \\ &= \max_{y_2, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots \max_{y_1} P(y_2|y_1)P(x_2|y_2) \text{score}_1(y_1) \end{aligned}$$

Abstract away the score for all decisions till here into **score**


$$\text{score}_1(s) = P(s)P(x_1|s)$$



# Deriving the recursive algorithm

$$P(x_1, x_2, \dots, x_n, y_1, y_2, \dots, y_n) = P(y_1) \prod_{i=1}^{n-1} P(y_{i+1}|y_i) \prod_{i=1}^n P(x_i|y_i)$$

$$\begin{aligned} & \max_{y_1, y_2, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots P(y_2|y_1)P(x_2|y_2)P(y_1)P(x_1|y_1) \\ &= \max_{y_2, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots \max_{y_1} P(y_2|y_1)P(x_2|y_2)P(y_1)P(x_1|y_1) \\ &= \max_{y_2, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots \max_{y_1} P(y_2|y_1)P(x_2|y_2) \text{score}_1(y_1) \\ &= \max_{y_3, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots \max_{y_2} P(y_3|y_2)P(x_3|y_3) \max_{y_1} P(y_2|y_1)P(x_2|y_2) \text{score}_1(y_1) \end{aligned}$$

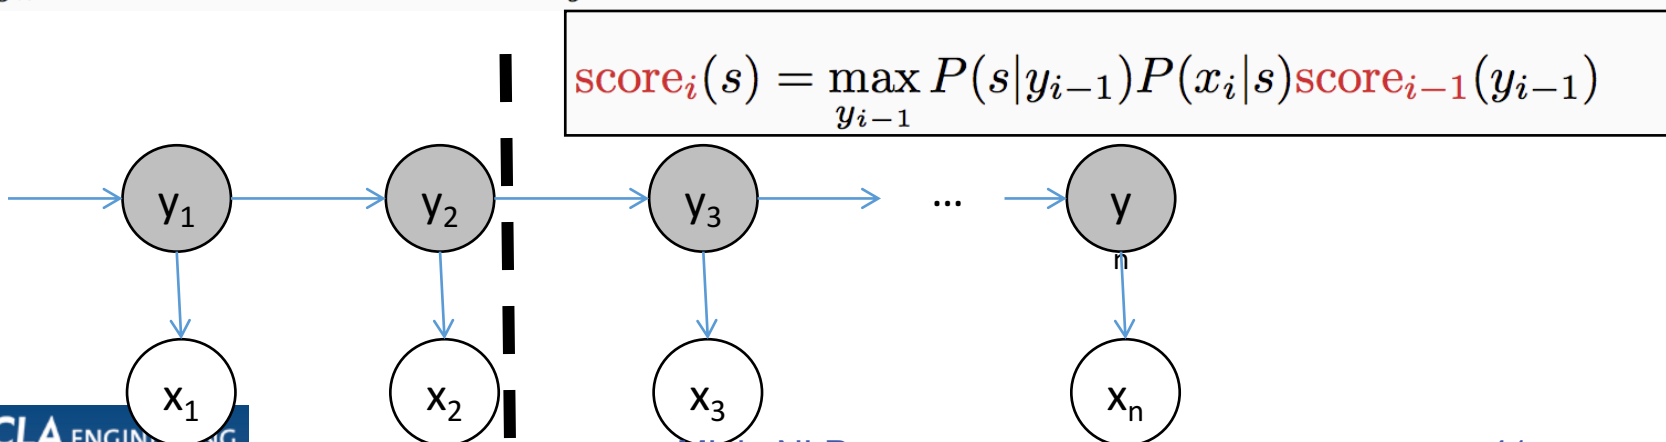




# Deriving the recursive algorithm

$$P(x_1, x_2, \dots, x_n, y_1, y_2, \dots, y_n) = P(y_1) \prod_{i=1}^{n-1} P(y_{i+1}|y_i) \prod_{i=1}^n P(x_i|y_i)$$

$$\begin{aligned} & \max_{y_1, y_2, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots P(y_2|y_1)P(x_2|y_2)P(y_1)P(x_1|y_1) \\ &= \max_{y_2, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots \max_{y_1} P(y_2|y_1)P(x_2|y_2)P(y_1)P(x_1|y_1) \\ &= \max_{y_2, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots \max_{y_1} P(y_2|y_1)P(x_2|y_2) \text{score}_1(y_1) \\ &= \max_{y_3, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots \max_{y_2} P(y_3|y_2)P(x_3|y_3) \max_{y_1} P(y_2|y_1)P(x_2|y_2) \text{score}_1(y_1) \\ &= \max_{y_3, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots \max_{y_2} P(y_3|y_2)P(x_3|y_3) \text{score}_2(y_2) \end{aligned}$$



Abstract away the score for all decisions till here into **score**

# Deriving the recursive algorithm

$$P(x_1, x_2, \dots, x_n, y_1, y_2, \dots, y_n) = P(y_1) \prod_{i=1}^{n-1} P(y_{i+1}|y_i) \prod_{i=1}^n P(x_i|y_i)$$

$$\max_{y_1, y_2, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots P(y_2|y_1)P(x_2|y_2)P(y_1)P(x_1|y_1)$$

$$= \max_{y_2, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots \max_{y_1} P(y_2|y_1)P(x_2|y_2)P(y_1)P(x_1|y_1)$$


$$= \max_{y_2, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots \max_{y_1} P(y_2|y_1)P(x_2|y_2) \text{score}_1(y_1)$$

$$= \max_{y_3, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots \max_{y_2} P(y_3|y_2)P(x_3|y_3) \max_{y_1} P(y_2|y_1)P(x_2|y_2) \text{score}_1(y_1)$$

$$= \max_{y_3, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots \max_{y_2} P(y_3|y_2)P(x_3|y_3) \text{score}_2(y_2)$$

⋮

$$= \max_{y_n} \text{score}_n(y_n)$$



# Deriving the recursive algorithm

$$P(x_1, x_2, \dots, x_n, y_1, y_2, \dots, y_n) = P(y_1) \prod_{i=1}^{n-1} P(y_{i+1}|y_i) \prod_{i=1}^n P(x_i|y_i)$$

$$\begin{aligned} & \max_{y_1, y_2, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots P(y_2|y_1)P(x_2|y_2)P(y_1)P(x_1|y_1) \\ &= \max_{y_2, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots \max_{y_1} P(y_2|y_1)P(x_2|y_2)P(y_1)P(x_1|y_1) \\ &= \max_{y_2, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots \max_{y_1} P(y_2|y_1)P(x_2|y_2) \text{score}_1(y_1) \\ &= \max_{y_3, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots \max_{y_2} P(y_3|y_2)P(x_3|y_3) \max_{y_1} P(y_2|y_1)P(x_2|y_2) \text{score}_1(y_1) \\ &= \max_{y_3, \dots, y_n} P(y_n|y_{n-1})P(x_n|y_n) \cdots \max_{y_2} P(y_3|y_2)P(x_3|y_3) \text{score}_2(y_2) \\ &\vdots \\ &= \max_{y_n} \text{score}_n(y_n) \end{aligned} \quad \text{score}_1(s) = P(s)P(x_1|s)$$

$$\text{score}_i(s) = \max_{y_{i-1}} P(s|y_{i-1})P(x_i|s) \text{score}_{i-1}(y_{i-1})$$

# Viterbi algorithm

Max-product algorithm for first order sequences

$\pi$ : Initial probabilities

A: Transitions

B: Emissions

1. **Initial**: For each state  $s$ , calculate

$$\text{score}_1(s) = P(s)P(x_1|s) = \pi_s B_{x_1,s}$$

1. **Recurrence**: For  $i = 2$  to  $n$ , for every state  $s$ , calculate

$$\begin{aligned}\text{score}_i(s) &= \max_{y_{i-1}} P(s|y_{i-1})P(x_i|s)\text{score}_{i-1}(y_{i-1}) \\ &= \max_{y_{i-1}} A_{y_{i-1},s} B_{s,x_i} \text{score}_{i-1}(y_{i-1})\end{aligned}$$

1. **Final state**: calculate

$$\max_{\mathbf{y}} P(\mathbf{y}, \mathbf{x}|\pi, A, B) = \max_s \text{score}_n(s)$$

This only calculates the max. To get final answer (*argmax*),

- keep track of which state corresponds to the max at each step
- build the answer using these back pointers

# General idea

- ❖ Dynamic programming

- ❖ The best solution for the full problem relies on best solution to sub-problems
- ❖ Memoize partial computation

- ❖ Examples

- ❖ Viterbi algorithm
- ❖ Dijkstra's shortest path algorithm
- ❖ ...

# Complexity of inference

- ❖ Complexity parameters

- ❖ Input sequence length:  $n$

- ❖ Number of states:  $K$

- ❖ Memory

- ❖ Storing the table:  $nK$  (scores for all states at each position)

- ❖ Runtime

- ❖ At each step, go over pairs of states

- ❖  $O(nK^2)$

# Outline

- ❖ Sequence models
- ❖ Hidden Markov models
  - ❖ Inference with HMM
  - ❖ *Learning*
- ❖ Conditional Models and Local Classifiers
- ❖ Global models
  - ❖ Conditional Random Fields
  - ❖ Structured Perceptron for sequences

# Learning HMM parameters

- ❖ Assume we know the number of states in the HMM
- ❖ Two possible scenarios

- 
1. We are given a data set  $D = \{<\mathbf{x}_i, \mathbf{y}_i>\}$  of sequences labeled with states

And we have to learn the parameters of the HMM  $(\pi, A, B)$

Supervised learning with complete data

2. We are given only a collection of sequences  $D = \{\mathbf{x}_i\}$

And we have to learn the parameters of the HMM  $(\pi, A, B)$

Unsupervised learning, with incomplete data



# Supervised learning of HMM


- ❖ We are given a dataset  $D = \{<\mathbf{x}_i, \mathbf{y}_i>\}$ 
  - ❖ Each  $\mathbf{x}_i$  is a sequence of observations and  $\mathbf{y}_i$  is a sequence of states that correspond to  $\mathbf{x}_i$

**Goal:** Learn initial, transition, emission distributions  $(\pi, A, B)$

- ❖ How do we learn the parameters of the probability distribution?

- ❖ **The maximum likelihood principle**

Where have we seen this before?

$$(\hat{\pi}, \hat{A}, \hat{B}) = \max_{\pi, A, B} P(D|\pi, A, B) = \max_{\pi, A, B} \prod_i P(\mathbf{x}_i, \mathbf{y}_i | \pi, A, B)$$


And we know how to write this in terms of the parameters of the HMM

# Supervised learning details

$$(\hat{\pi}, \hat{A}, \hat{B}) = \max_{\pi, A, B} P(D|\pi, A, B) = \max_{\pi, A, B} \prod_i P(\mathbf{x}_i, \mathbf{y}_i|\pi, A, B)$$

$\pi, A, B$  can be estimated separately just by counting

❖ Makes learning simple and fast

[**Exercise**: Derive the following using derivatives of the log likelihood.  
Requires Lagrangian multipliers.]

$$\pi_s = \frac{\text{count}(\text{start} \rightarrow s)}{n}$$

Initial probabilities

Number of instances where the first state is  $s$

Number of examples

$$A_{s',s} = \frac{\text{count}(s \rightarrow s')}{\text{count}(s)}$$

Transition probabilities

$$B_{s,x} = \frac{\text{count} \left( \begin{array}{c} s \\ \downarrow \\ x \end{array} \right)}{\text{count}(s)}$$

Emission probabilities

# Hidden Markov Models summary

- ❖ Predicting sequences
  - ❖ As many output states as observations
- ❖ Markov assumption helps decompose the score
- ❖ Several algorithmic questions
  - ❖ Most likely state
  - ❖ Learning parameters
    - ❖ Supervised, Unsupervised
  - ❖ Probability of an observation sequence
    - ❖ Sum over all assignments to states, replace max with sum in Viterbi
  - ❖ Probability of state for each observation
    - ❖ Sum over all assignments to all other states

# Outline

- ❖ Sequence models
- ❖ Hidden Markov models
  - ❖ Inference with HMM
  - ❖ Learning
- ❖ Conditional Models and Local Classifiers
- ❖ Global models
  - ❖ Conditional Random Fields

# Outline

- ❖ Sequence models
- ❖ Hidden Markov models
  - ❖ Inference with HMM
  - ❖ Learning
- ❖ Conditional Models and Local Classifiers
- ❖ Global models
  - ❖ Conditional Random Fields

# Modeling next-state directly

- ❖ Instead of modeling the joint distribution  $P(\mathbf{x}, \mathbf{y})$  only focus on  $P(\mathbf{y}|\mathbf{x})$ 
  - ❖ Which is what we care about eventually anyway
- ❖ For sequences, different formulations
  - ❖ Maximum Entropy Markov Model [McCallum, et al 2000]
  - ❖ Projection-based Markov Model [Punyakanok and Roth, 2001]  
(other names: discriminative/conditional markov model, ...)

# Generative vs Discriminative models

## ❖ Generative models

- ❖ learn  $P(x, y)$
- ❖ Characterize how the data is generated (both inputs and outputs)
- ❖ Eg: Naïve Bayes, Hidden Markov Model

## ❖ Discriminative models

- ❖ learn  $P(y | x)$
- ❖ Directly characterizes the decision boundary only
- ❖ Eg: Logistic Regression, Conditional models (several names)

# Generative vs Discriminative models


## ❖ Generative models

- ❖ learn  $P(x, y)$
- ❖ Characterize how the data is generated (both **inputs** and **outputs**)
- ❖ Eg: Naïve Bayes, Hidden Markov Model

A generative model tries to characterize the distribution of the inputs, a discriminative model doesn't care

## ❖ Discriminative models



- ❖ learn  $P(y | x)$
- ❖ Directly characterizes the decision **boundary** only
- ❖ Eg: Logistic Regression, Conditional models (several names)





# Another independence assumption

$$P(y_i | \textcolor{red}{y}_{i-1}, y_{i-2}, \dots, \textcolor{red}{x}_i, x_{i-1}, \dots) = P(y_i | y_{i-1}, x_i)$$



This assumption lets us write the conditional probability of the output as

$$P(\mathbf{y} | \mathbf{x}) = \prod_i P(y_i | y_{i-1}, x_i)$$

# Modeling $P(y_i \mid y_{i-1}, x_i)$

## ❖ Different approaches possible

1. Train a *log-linear* classifier
2. Or, ignore the fact that we are predicting a probability, we only care about maximizing some *score*. Train any classifier (e.g, perceptron algorithm)

## ❖ For both cases:

- ❖ Use rich features that depend on input and previous state
- ❖ We can increase the dependency to arbitrary neighboring  $x_i$ 's
  - ❖ Eg. Neighboring words influence this words POS tag

# Log-linear models for multiclass

## Consider multiclass classification

- ❖ Inputs:  $\mathbf{x}$
- ❖ Output:  $\mathbf{y} \in \{1, 2, \dots, K\}$
- ❖ Feature representation:  $\phi(\mathbf{x}, \mathbf{y})$ 
  - ❖ We have seen this before
- ❖ Define probability of an input  $\mathbf{x}$  taking a label  $\mathbf{y}$  as

$$P(\mathbf{y}|\mathbf{x}, \mathbf{w}) = \frac{e^{\mathbf{w}^T \phi(\mathbf{x}, \mathbf{y})}}{\sum_{y'} e^{\mathbf{w}^T \phi(\mathbf{x}, y')}}}$$

**Interpretation:** Score for label, converted to a well-formed probability distribution by exponentiating + normalizing

- ❖ A generalization of logistic regression to multiclass

# Training a log-linear model

Given a data set  $D = \{ \langle \mathbf{x}_i, \mathbf{y}_i \rangle \}$

❖ Apply the **maximum likelihood** principle

$$\max_{\mathbf{w}} \sum_i \log P(\mathbf{y}_i | \mathbf{x}_i, \mathbf{w})$$

❖ Maybe with a **regularizer**

Here

$$P(\mathbf{y} | \mathbf{x}, \mathbf{w}) = \frac{e^{\mathbf{w}^T \phi(\mathbf{x}, \mathbf{y})}}{\sum_{\mathbf{y}'} e^{\mathbf{w}^T \phi(\mathbf{x}, \mathbf{y}')}}$$

$$\max_{\mathbf{w}} -\frac{\lambda}{2} \mathbf{w}^T \mathbf{w} + \sum_i \log P(\mathbf{y}_i | \mathbf{x}_i, \mathbf{w})$$

# Training a log-linear model



- ❖ Gradient based methods to minimize

$$L(\mathbf{w}) = \sum_i \log P(\mathbf{y}_i | \mathbf{x}_i, \mathbf{w})$$

- ❖ Usual stochastic gradient descent
  - ❖ Initialize  $\mathbf{w} \leftarrow \mathbf{0}$
  - ❖ Iterate through examples for multiple epochs
    - ❖ For each example  $(\mathbf{x}_i, \mathbf{y}_i)$  take gradient step for the loss at that example
      - ❖ Update  $\mathbf{w} \leftarrow \mathbf{w} - r_t \nabla L(\mathbf{w}, \mathbf{x}_i, \mathbf{y}_i)$
  - ❖ Return  $\mathbf{w}$

# The next-state model

$$P(y_i | \textcolor{red}{y}_{i-1}, y_{i-2}, \dots, \textcolor{red}{x}_i, x_{i-1}, \dots) = P(y_i | y_{i-1}, x_i)$$



This assumption lets us write the conditional probability of the output as

$$P(\mathbf{y}|\mathbf{x}) = \prod_i \boxed{P(y_i | y_{i-1}, x_i)}$$

We need to learn this function

# Modeling ~~$P(y_i | y_{i-1}, x_i)$~~ $P(y_i | y_{i-1}, \mathbf{x})$

## ❖ Different approaches possible

1. Train a *maximum entropy* classifier

Basically, multinomial logistic regression

2. Ignore the fact that we are predicting a probability, we only care about maximizing some *score*. Train any classifier, using say the perceptron algorithm


## ❖ For both cases:

- ❖ Use rich features that depend on input and previous state
- ❖ We can increase the dependency to arbitrary neighboring  $x_i$ 's
  - ❖ Eg. Neighboring words influence this words POS tag

# Maximum Entropy Markov Model

Goal: Compute  $P(\mathbf{y} \mid \mathbf{x})$

$$P(y_i | y_{i-1}, \mathbf{x}) \propto \exp(\mathbf{w}^T \phi(\mathbf{x}, i, y_i, y_{i-1}))$$



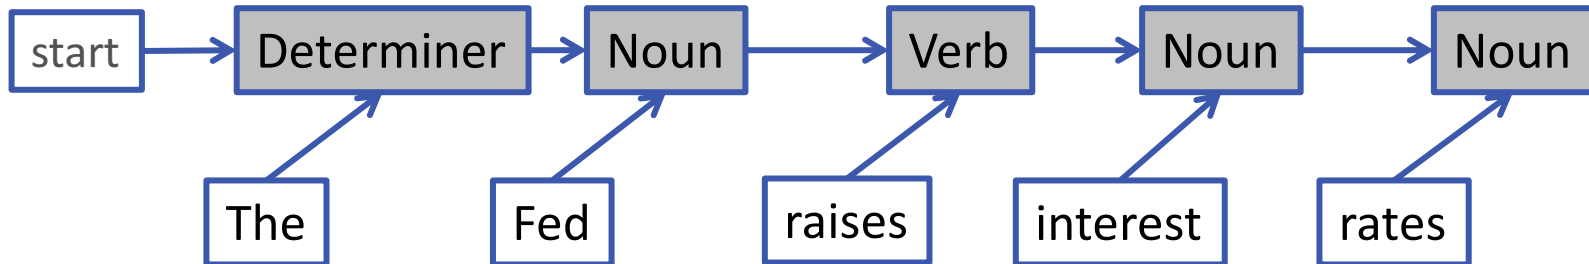
The prediction task:

Using the entire input and the current label, predict the next label



# Maximum Entropy Markov Model

Goal: Compute  $P(\mathbf{y} \mid \mathbf{x})$   $P(y_i | y_{i-1}, \mathbf{x}) \propto \exp(\mathbf{w}^T \phi(\mathbf{x}, i, y_i, y_{i-1}))$



*word*

*Caps*

*-es*

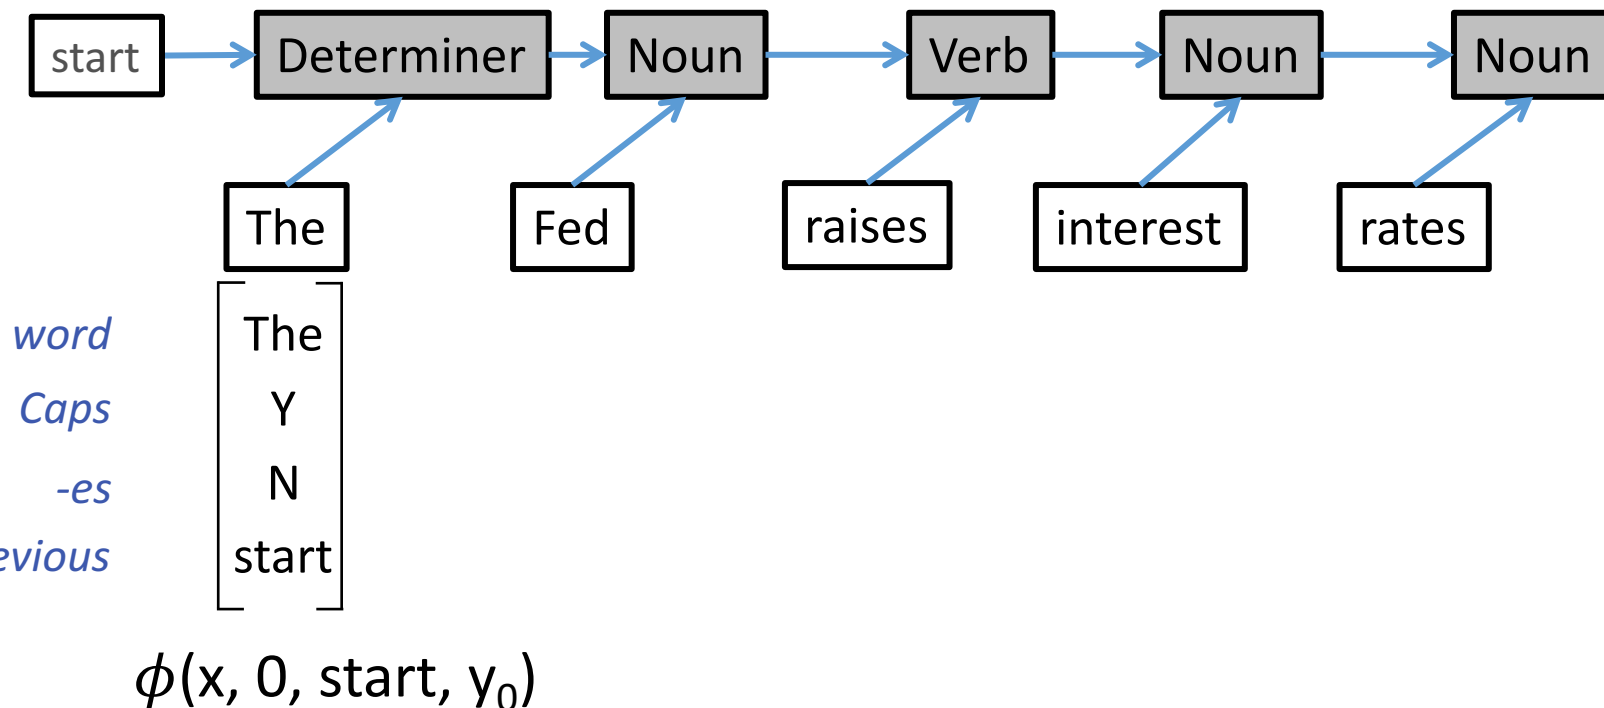
To model the probability, first, we need to define features for the current classification problem

*Previous*

# Maximum Entropy Markov Model

Goal: Compute  $P(\mathbf{y} \mid \mathbf{x})$


$$P(y_i | y_{i-1}, \mathbf{x}) \propto \exp(\mathbf{w}^T \phi(\mathbf{x}, i, y_i, y_{i-1}))$$



# Maximum Entropy Markov Model

Goal: Compute  $P(\mathbf{y} \mid \mathbf{x})$

$$P(y_i | y_{i-1}, \mathbf{x}) \propto \exp(\mathbf{w}^T \phi(\mathbf{x}, i, y_i, y_{i-1}))$$



*word*

*Caps*

*-es*

*Previous*

The  
Y  
N  
start

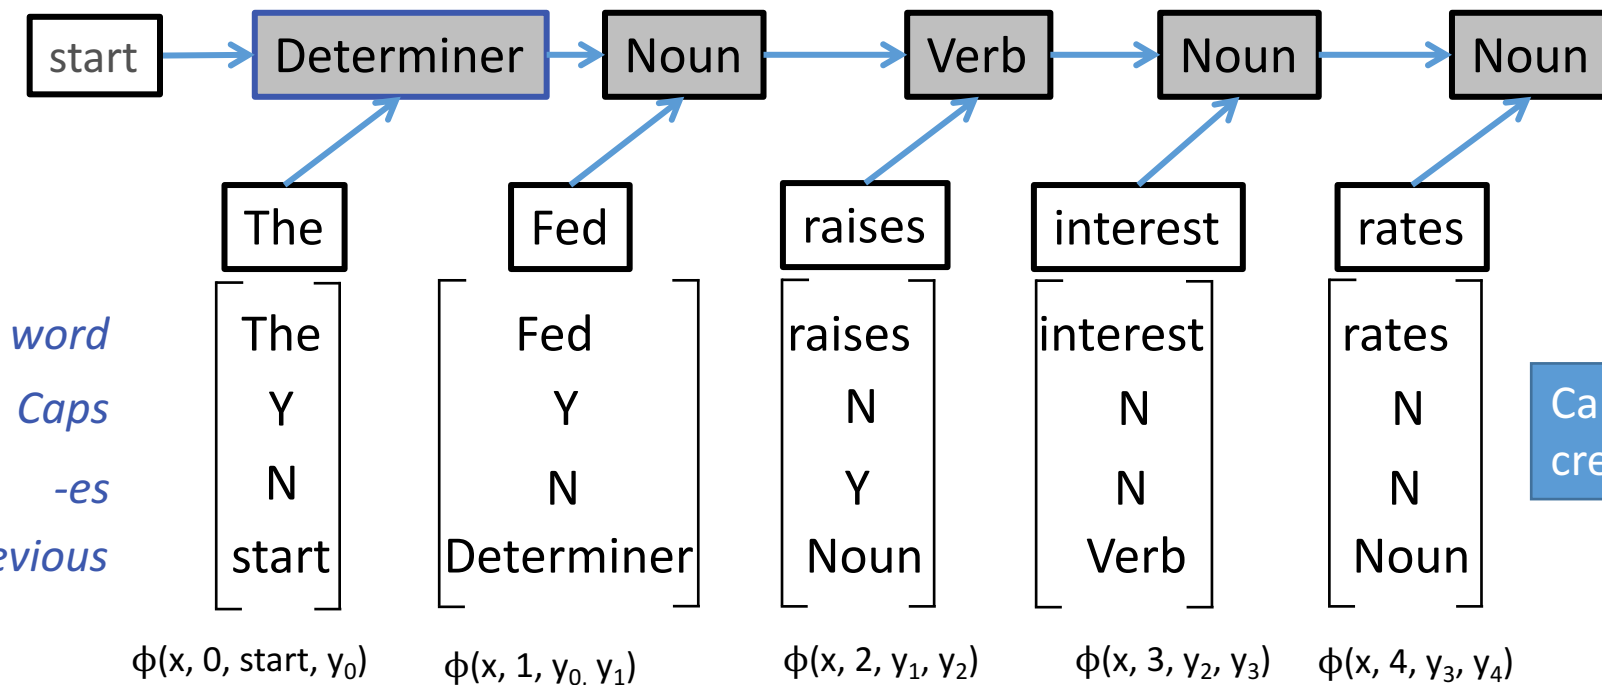
Fed  
Y  
N  
Determiner

$\phi(\mathbf{x}, 0, \text{start}, y_0)$   $\phi(\mathbf{x}, 1, y_0, y_1)$

# Maximum Entropy Markov Model

Goal: Compute  $P(\mathbf{y} \mid \mathbf{x})$

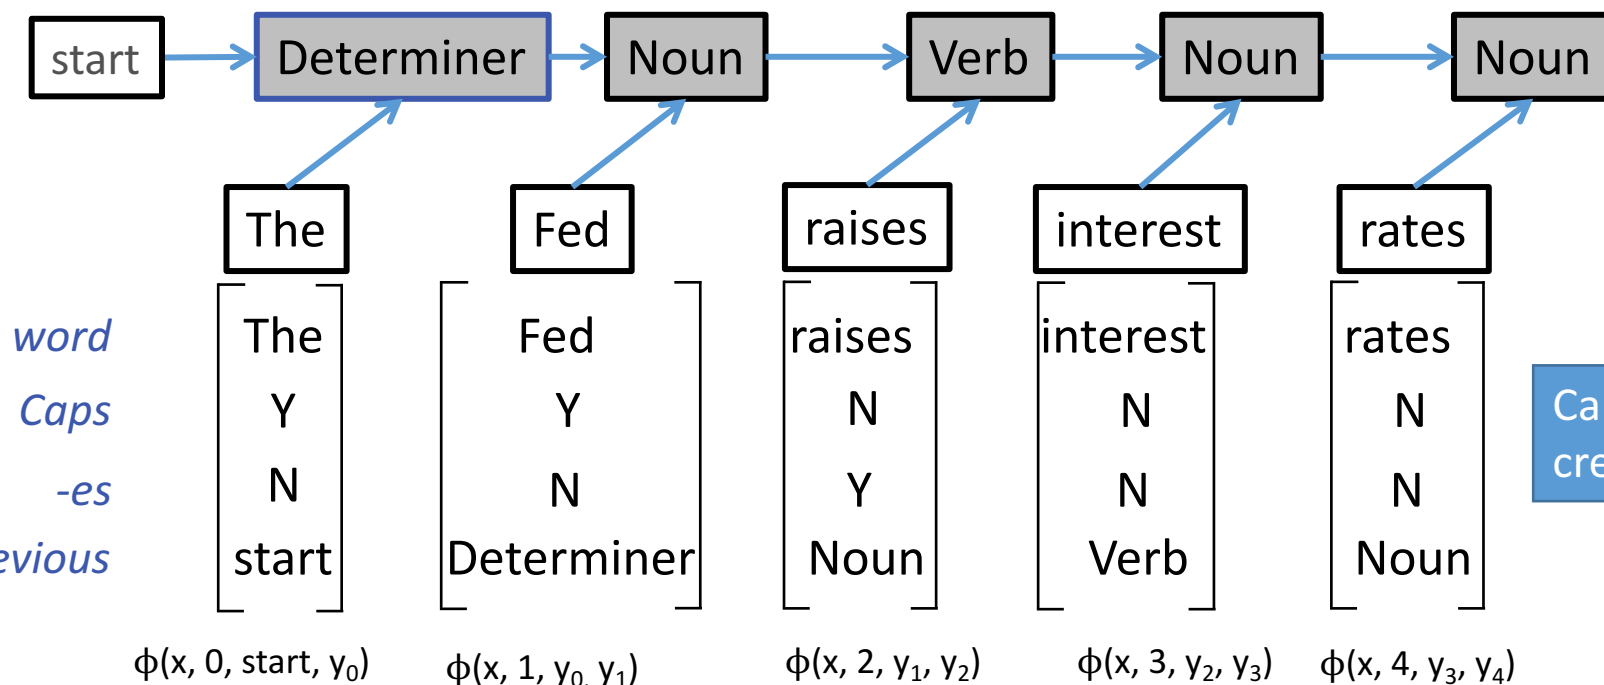
$$P(y_i | y_{i-1}, \mathbf{x}) \propto \exp(\mathbf{w}^T \phi(\mathbf{x}, i, y_i, y_{i-1}))$$




# Maximum Entropy Markov Model

Goal: Compute  $P(\mathbf{y} \mid \mathbf{x})$

$$P(y_i | y_{i-1}, \mathbf{x}) \propto \exp(\mathbf{w}^T \phi(\mathbf{x}, i, y_i, y_{i-1}))$$





Compare to HMM: Only depends on the word and the previous tag



# Using MEMM

- ❖ Training
  - ❖ Next-state predictor **locally** as maximum likelihood
    - ❖ Similar to any maximum entropy classifier
- ❖ Prediction/decoding
  - ❖ Modify the Viterbi algorithm for the new independence assumptions



$$\text{score}_i(s) = \max_{y_{i-1}} P(s|y_{i-1})P(x_i|s)\text{score}_{i-1}(y_{i-1})$$

$$\text{score}_i(s) = \max_{y_{i-1}} P(s|y_{i-1}, x_i)\text{score}_{i-1}(y_{i-1})$$

# Generalization: Any multiclass classifier

- ❖ **Viterbi decoding**: we only need a score for each decision
  - ❖ So far, probabilistic classifiers
- ❖ In general, use any learning algorithm to build get a score for the label  $y_i$  given  $y_{i-1}$  and  $\mathbf{x}$ 
  - ❖ Multiclass versions of perceptron, SVM
  - ❖ Just like MEMM, these allow arbitrary features to be defined

**Exercise:** Viterbi needs to be re-defined to work with sum of scores rather than the product of probabilities

# Comparison to HMM

## What we gain

### 1. Rich feature representation for inputs

- ❖ Helps generalize better by thinking about properties of the input tokens rather than the entire tokens
- ❖ Eg: If a word ends with –es, it might be a present tense verb (such as raises). Could be a feature; HMM cannot capture this

### 2. Discriminative predictor

- ❖ Model  $P(\mathbf{y} \mid \mathbf{x})$  rather than  $P(\mathbf{y}, \mathbf{x})$
- ❖ *Joint vs conditional*



# Outline



- ❖ Sequence models
- ❖ Hidden Markov models
  - ❖ Inference with HMM
  - ❖ Learning
- ❖ Conditional Models and Local Classifiers
- ❖ Global models
  - ❖ Conditional Random Fields
  - ❖ Structured Perceptron for sequences

# Outline

- ❖ Conditional models for predicting sequences
- ❖ Log-linear models for multiclass classification
- ❖ Maximum Entropy Markov Models
  - ❖ The Label Bias Problem

# The next-state model for sequences

$$P(y_i | \textcolor{red}{y}_{i-1}, y_{i-2}, \dots, \textcolor{red}{x}_i, x_{i-1}, \dots) = P(y_i | y_{i-1}, x_i)$$



This assumption lets us write the conditional probability of the output as

$$P(\mathbf{y}|\mathbf{x}) = \prod_i P(y_i | y_{i-1}, x_i)$$

We need to train local multiclass classifiers that predicts the next state given the previous state and the input

## ...local classifiers! Label bias problem

Let's look at the independence assumption

$$P(y_i | \mathbf{y}_{i-1}, y_{i-2}, \dots, \mathbf{x}_i, x_{i-1}, \dots) = P(y_i | y_{i-1}, x_i)$$

“Next-state” classifiers  
are locally normalized


# ...local classifiers! Label bias problem

Let's look at the independence assumption

$$P(y_i | \mathbf{y}_{i-1}, y_{i-2}, \dots, \mathbf{x}_i, x_{i-1}, \dots) = P(y_i | y_{i-1}, x_i)$$

“Next-state” classifiers are locally normalized


Eg: Part-of-speech tagging the sentence



Suppose these are the only state transitions allowed

# ...local classifiers → Label bias problem

*The robot wheels are round*



Suppose these are the only state transitions allowed

Option 1:  $P(D \mid \text{The}) \cdot$


$$P(N \mid D, \text{robot}) \cdot \\ P(N \mid N, \text{wheels}) \cdot \\ P(V \mid N, \text{are}) \cdot \\ P(A \mid V, \text{round})$$

Option 2:  $P(D \mid \text{The}) \cdot$

$$P(N \mid D, \text{robot}) \cdot \\ P(V \mid N, \text{wheels}) \cdot \\ P(N \mid V, \text{are}) \cdot \\ P(R \mid N, \text{round})$$

## ...local classifiers → Label bias problem

*The robot wheels are round*



Suppose these are the only state transitions allowed

Option 1:  $P(D \mid \text{The}) \cdot$

$P(N \mid D, \text{robot}) \cdot$

$P(N \mid N, \text{wheels}) \cdot$

~~$P(V \mid N, \text{are})$~~   $P(V \mid N, \text{Fred}) \cdot$

$P(A \mid V, \text{round})$

Option 2:  $P(D \mid \text{The}) \cdot$

$P(N \mid D, \text{robot}) \cdot$

$P(V \mid N, \text{wheels}) \cdot$

~~$P(N \mid V, \text{are})$~~   $P(N \mid V, \text{Fred}) \cdot$

$P(R \mid N, \text{round})$

The path scores are the same

Even if the word Fred is never observed as a verb in the data,  
it will be predicted as one

The input Fred does not influence the output at all

# Label Bias

- ❖ States with a single outgoing transition effectively ignore their input
  - ❖ States with lower-entropy next states are less influenced by observations
- ❖ Why?
  - ❖ Each the next-state classifiers are locally normalized.
  - ❖ If a state has fewer next states, each of those will get a higher probability mass  
...and hence preferred
- ❖ Side note: Surprisingly doesn't affect some tasks
  - ❖ Eg: part-of-speech tagging



# Summary: Local models for Sequences

- ❖ Conditional models
- ❖ Use rich features in the mode
- ❖ Possibly suffer from label bias problem

# Outline

- ❖ Sequence models
- ❖ Hidden Markov models
  - ❖ Inference with HMM
  - ❖ Learning
- ❖ Conditional Models and Local Classifiers
- ❖ Global models
  - ❖ Conditional Random Fields

# Outline

- ❖ Sequence models
- ❖ Hidden Markov models
  - ❖ Inference with HMM
  - ❖ Learning
- ❖ Conditional Models and Local Classifiers
- ❖ Global models
  - ❖ *Conditional Random Fields*

# So far...

## ❖ Hidden Markov models

- ❖ **Pros:** Decomposition of total probability with tractable
- ❖ **Cons:** Doesn't allow use of features for representing inputs
  - ❖ Also, generative model
    - ❖ not really a downside, but we may get better performance with conditional models if we care only about predictions

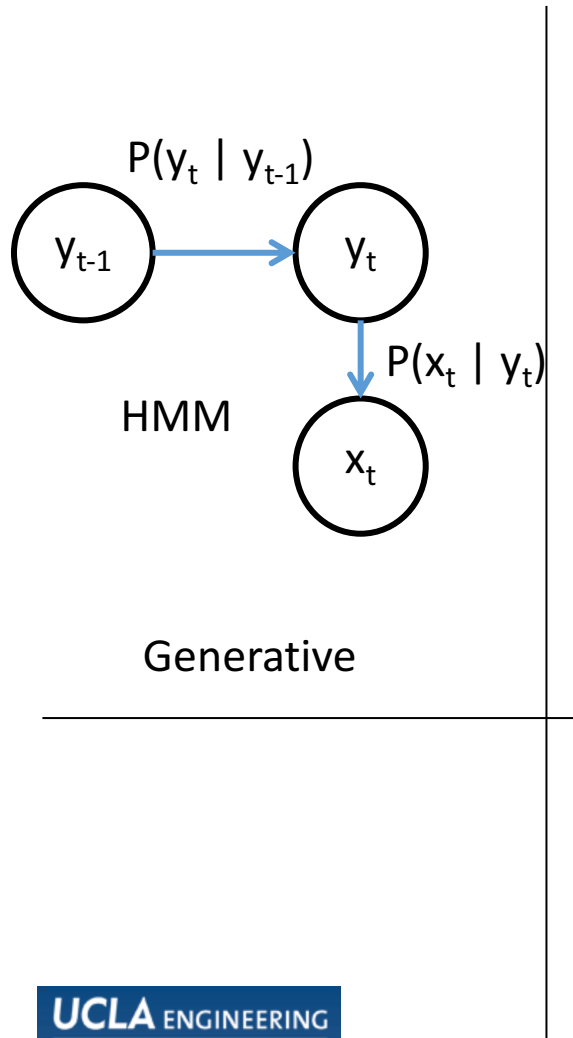
## ❖ Local, conditional Markov Models

- ❖ **Pros:** Conditional model, allows features to be used
- ❖ **Cons:** Label bias problem

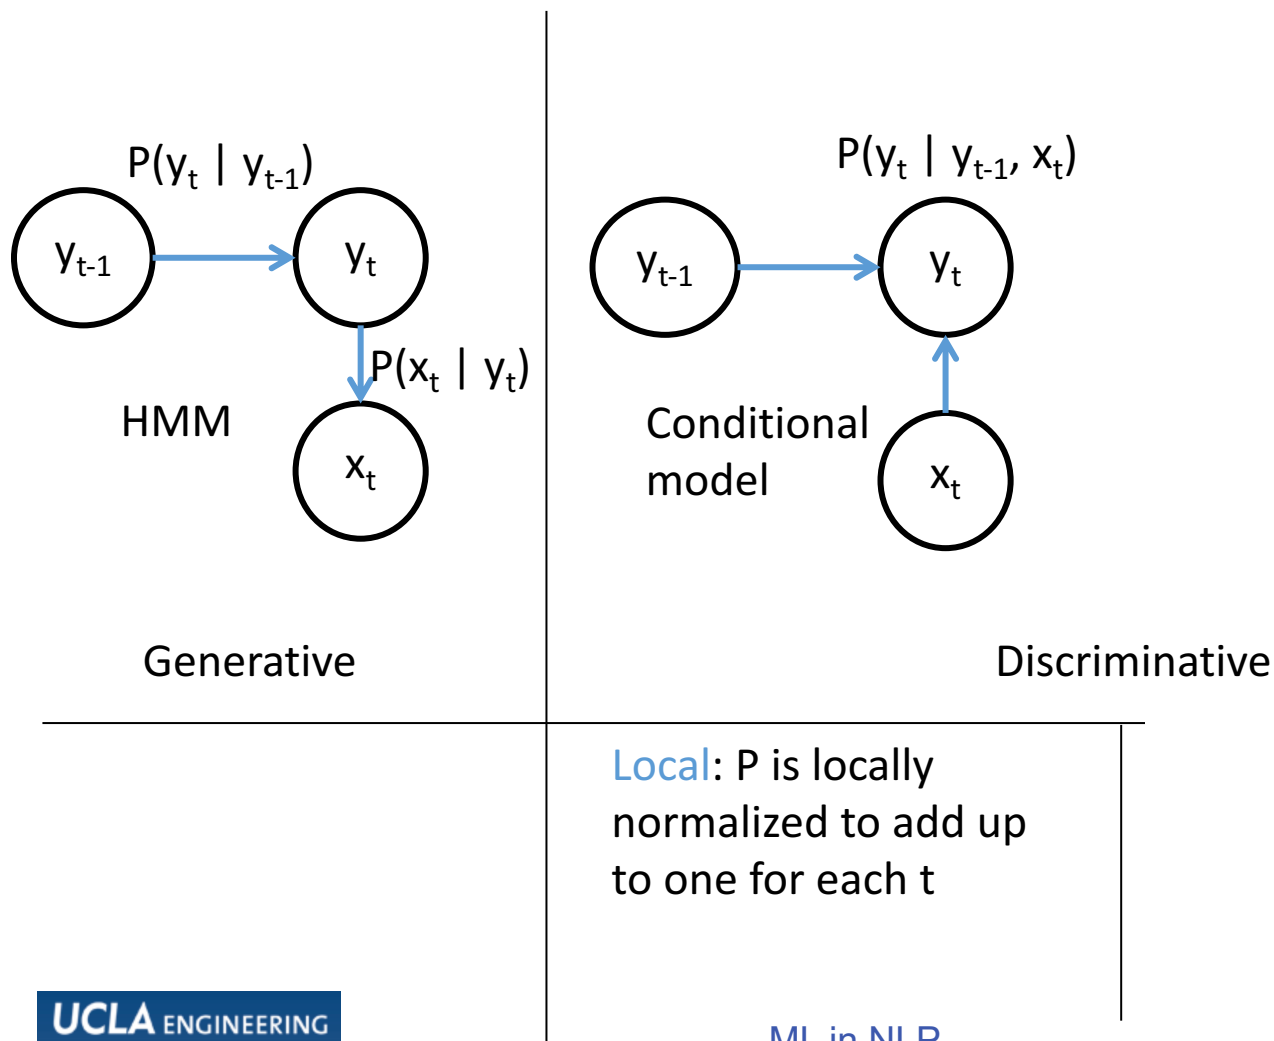
# Global models

- ❖ Train the predictor globally
  - ❖ Instead of training local decisions independently
- ❖ Normalize globally
  - ❖ Make each edge in the model undirected
  - ❖ Not associated with a probability, but just a “score”
- ❖ Recall the difference between local vs. global for multiclass

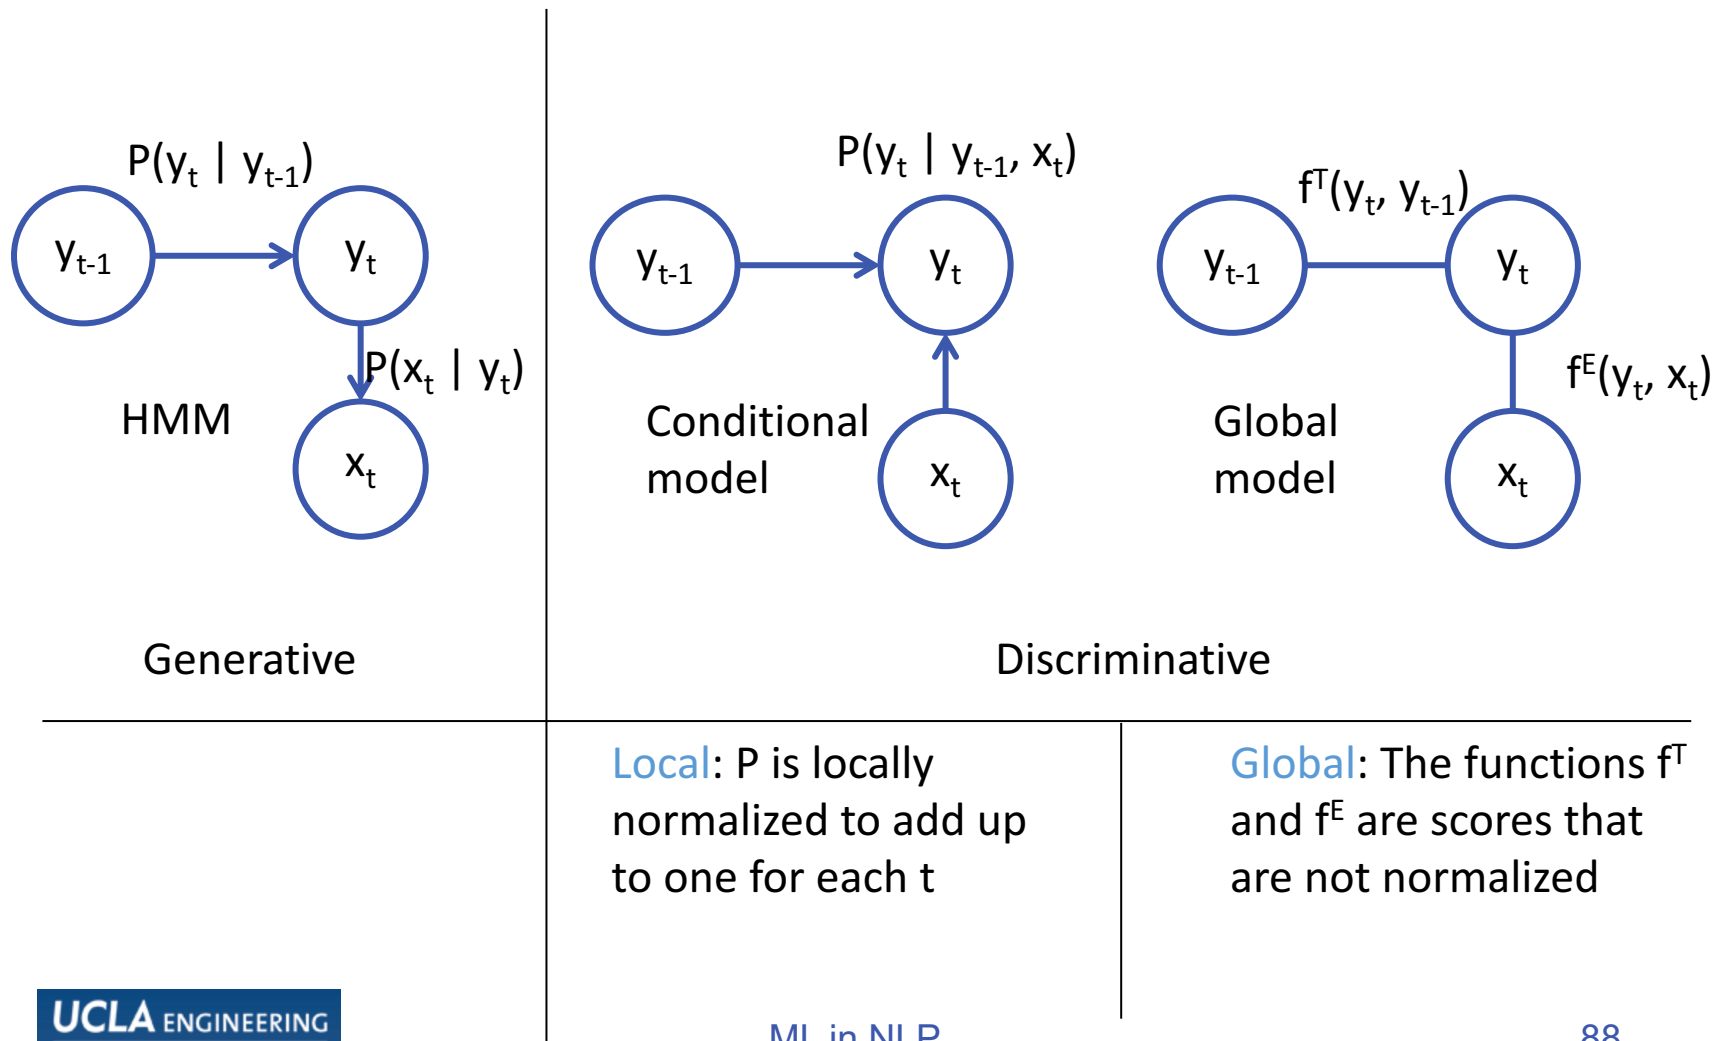
# HMM vs. A local model vs. A global model



# HMM vs. A local model vs. A global model




# HMM vs. A local model vs. A global model





# Conditional Random Field




Each node is a random variable

We observe some nodes and the rest are unobserved

**The goal:** To characterize a probability distribution over the unobserved variables, given the observed

# Conditional Random Field




Each node is a random variable

We observe some nodes and need to assign the rest

Each **clique** is associated with a score

# Conditional Random Field

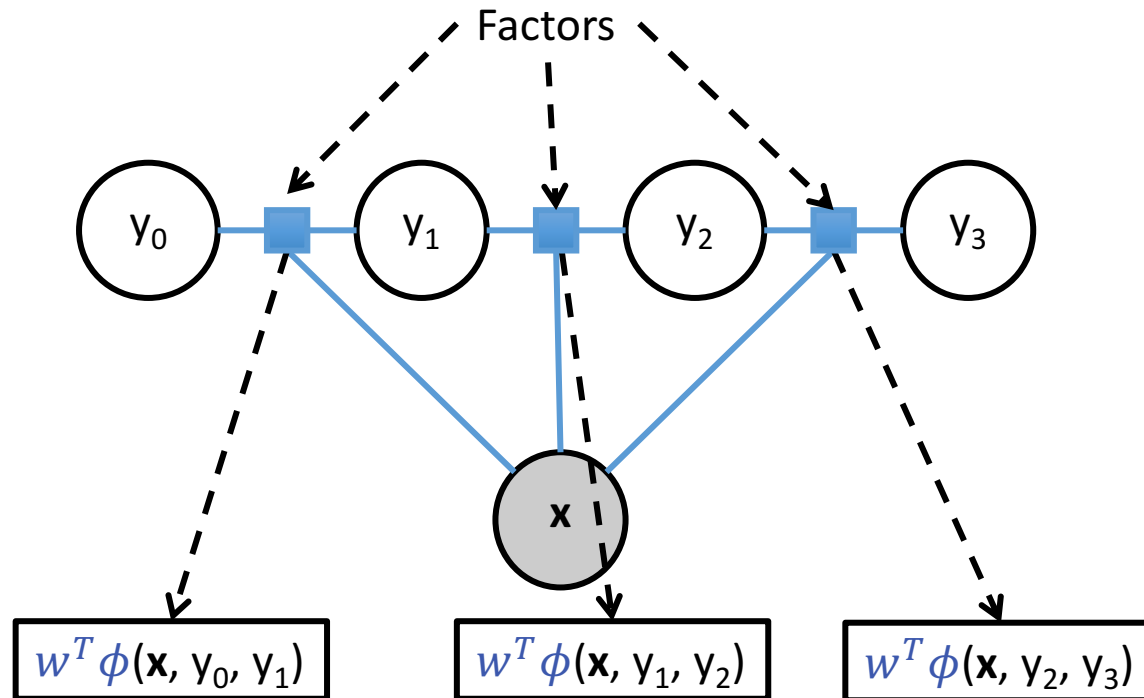


Each node is a random variable

We observe some nodes and need to assign the rest

Each **clique** is associated with a score

# Conditional Random Field: Factor graph



Each node is a random variable

We observe some nodes and need to assign the rest


Each **clique** is associated with a score

**factor**

# Conditional Random Field: Factor graph

A different factorization:

Recall decomposition of structures into parts. Same idea




Each node is a random variable

We observe some nodes and need to assign the rest

Each factor is associated with a score

# Conditional Random Field: Factor graph



Each node is a random variable

We observe some nodes and need to assign the rest


Each factor is associated with a score

# Conditional Random Field for sequences

$$P(\mathbf{y} \mid \mathbf{x}) = \frac{1}{Z} \prod_i \exp(\mathbf{w}^T \phi(\mathbf{x}, y_i, y_{i-1}))$$

Z: Normalizing constant,  
sum over all sequences

$$Z = \sum_{\hat{\mathbf{y}}} \prod_i \exp(w^T \phi(\mathbf{x}, \hat{y}_i, \hat{y}_{i-1}))$$



# CRF: A different view

- ❖ Input:  $\mathbf{x}$ , Output:  $\mathbf{y}$ , both sequences (for now)
- ❖ Define a feature vector for the **entire** input and output sequence:  $\phi(\mathbf{x}, \mathbf{y})$

- ❖ Define a giant log-linear model,  $P(\mathbf{y} \mid \mathbf{x})$  parameterized by  $\mathbf{w}$

$$P(\mathbf{y} \mid \mathbf{x}) = \frac{1}{Z} \prod_i \exp(\mathbf{w}^T \phi(\mathbf{x}, y_i, y_{i-1})) \propto \exp\left(\mathbf{w}^T \sum_i \phi(\mathbf{x}, y_i, y_{i-1})\right)$$

- ❖ Just like any other log-linear model, except
  - ❖ Space of  $\mathbf{y}$  is the set of all possible sequences of the correct length
  - ❖ Normalization constant sums over all sequences

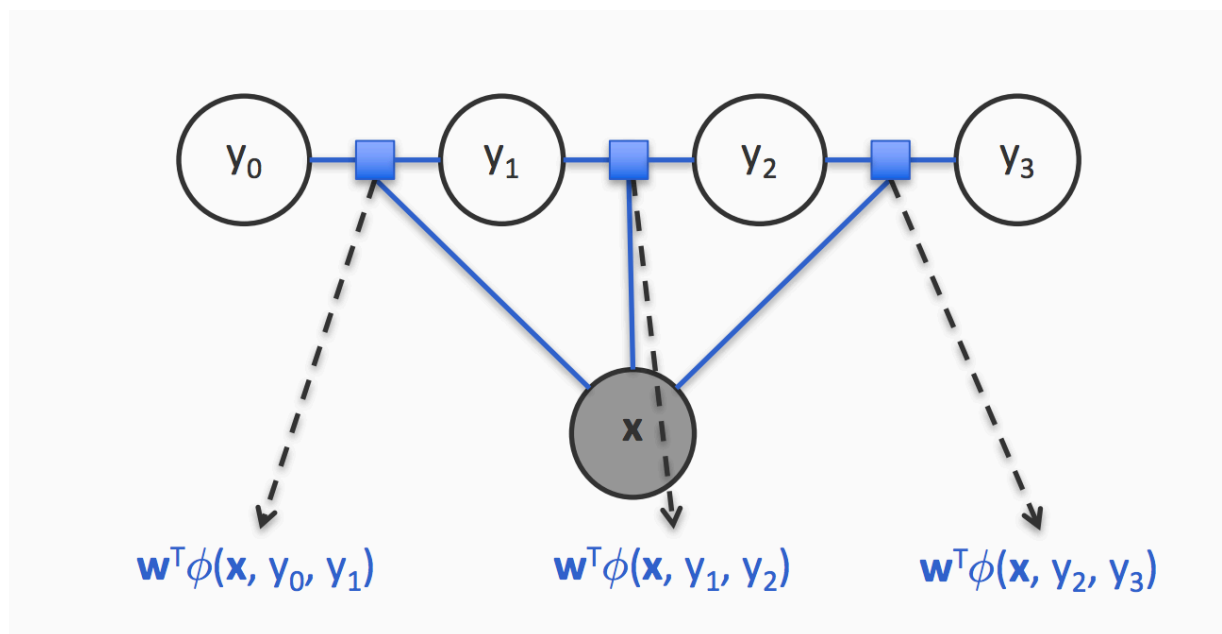
*In an MEMM, probabilities were locally normalized*



# Global features

The feature function decomposes over the sequence

$$\phi(\mathbf{x}, \mathbf{y}) = \sum_i \phi(\mathbf{x}, y_i, y_{i-1})$$



# Prediction

Goal: To predict most probable sequence  $\mathbf{y}$  an input  $\mathbf{x}$

$$\arg \max_{\mathbf{y}} P(\mathbf{y}|\mathbf{x}) = \arg \max_{\mathbf{y}} \exp(\mathbf{w}^T \phi(\mathbf{x}, \mathbf{y})) = \arg \max_{\mathbf{y}} \mathbf{w}^T \phi(\mathbf{x}, \mathbf{y})$$

But the score decomposes as  $\mathbf{w}^T \phi(\mathbf{x}, \mathbf{y}) = \sum_i \mathbf{w}^T \phi(\mathbf{x}, y_i, y_{i-1})$

Prediction via Viterbi (with sum instead of product)

$$\text{score}_0(s) = \mathbf{w}^T \phi(\mathbf{x}, y_0, \text{start})$$

$$\text{score}_i(s) = \max_{y_{i-1}} (\mathbf{w}^T \phi(\mathbf{x}, y_i, y_{i-1}) + \text{score}_{i-1}(y_{i-1}))$$

# Training a chain CRF

## ❖ Input:

- ❖ Dataset with labeled sequences,  $D = \{<\mathbf{x}_i, \mathbf{y}_i>\}$
- ❖ A definition of the feature function

## ❖ How do we train?

- ❖ Maximize the (regularized) log-likelihood

$$\max_{\mathbf{w}} -\frac{\lambda}{2} \mathbf{w}^T \mathbf{w} + \sum_i \log P(\mathbf{y}_i | \mathbf{x}_i, \mathbf{w})$$

Recall: Empirical loss minimization

# Training with inference

- ❖ Many methods for training  $\max_{\mathbf{w}} -\frac{\lambda}{2} \mathbf{w}^T \mathbf{w} + \sum_i \log P(\mathbf{y}_i | \mathbf{x}_i, \mathbf{w})$
- ❖ Numerical optimization
- ❖ Can use a gradient or hessian based method

- ❖ Simple gradient ascent

$$\mathbf{w} \leftarrow \mathbf{w} + \sum_i \left( \phi(\mathbf{x}_i, \mathbf{y}_i) - \sum_{\hat{\mathbf{y}}} P(\hat{\mathbf{y}} | \mathbf{x}_i, \mathbf{w}) \phi(\mathbf{x}_i, \hat{\mathbf{y}}) \right)$$

# Training with inference

- ❖ Many methods for training

- ❖ Numerical optimization

- ❖ Can use a gradient or hessian based method

$$\max_{\mathbf{w}} -\frac{\lambda}{2} \mathbf{w}^T \mathbf{w} + \sum_i \log P(\mathbf{y}_i | \mathbf{x}_i, \mathbf{w})$$

- ❖ Simple gradient ascent

$$\mathbf{w} \leftarrow \mathbf{w} + \sum_i \left( \phi(\mathbf{x}_i, \mathbf{y}_i) - \sum_{\hat{\mathbf{y}}} P(\hat{\mathbf{y}} | \mathbf{x}_i, \mathbf{w}) \phi(\mathbf{x}_i, \hat{\mathbf{y}}) \right)$$

- ❖ Training involves inference!

- ❖ A different kind than what we have seen so far

- ❖ Summing over all sequences is just like Viterbi

- ❖ With summation instead of maximization

# CRF summary

- ❖ An undirected graphical model
  - ❖ Decompose the score over the structure into a collection of factors
  - ❖ Each factor assigns a score to assignment of the random variables it is connected to
- ❖ Training and prediction
  - ❖ Final prediction via  $\operatorname{argmax} w^T \phi(\mathbf{x}, \mathbf{y})$
  - ❖ Train by maximum (regularized) likelihood (also need inference)
- ❖ Relation to other models
  - ❖ Effectively a linear classifier
  - ❖ A generalization of logistic regression to structures
  - ❖ An instance of Markov Random Field, with some random variables observed (We will see this soon)